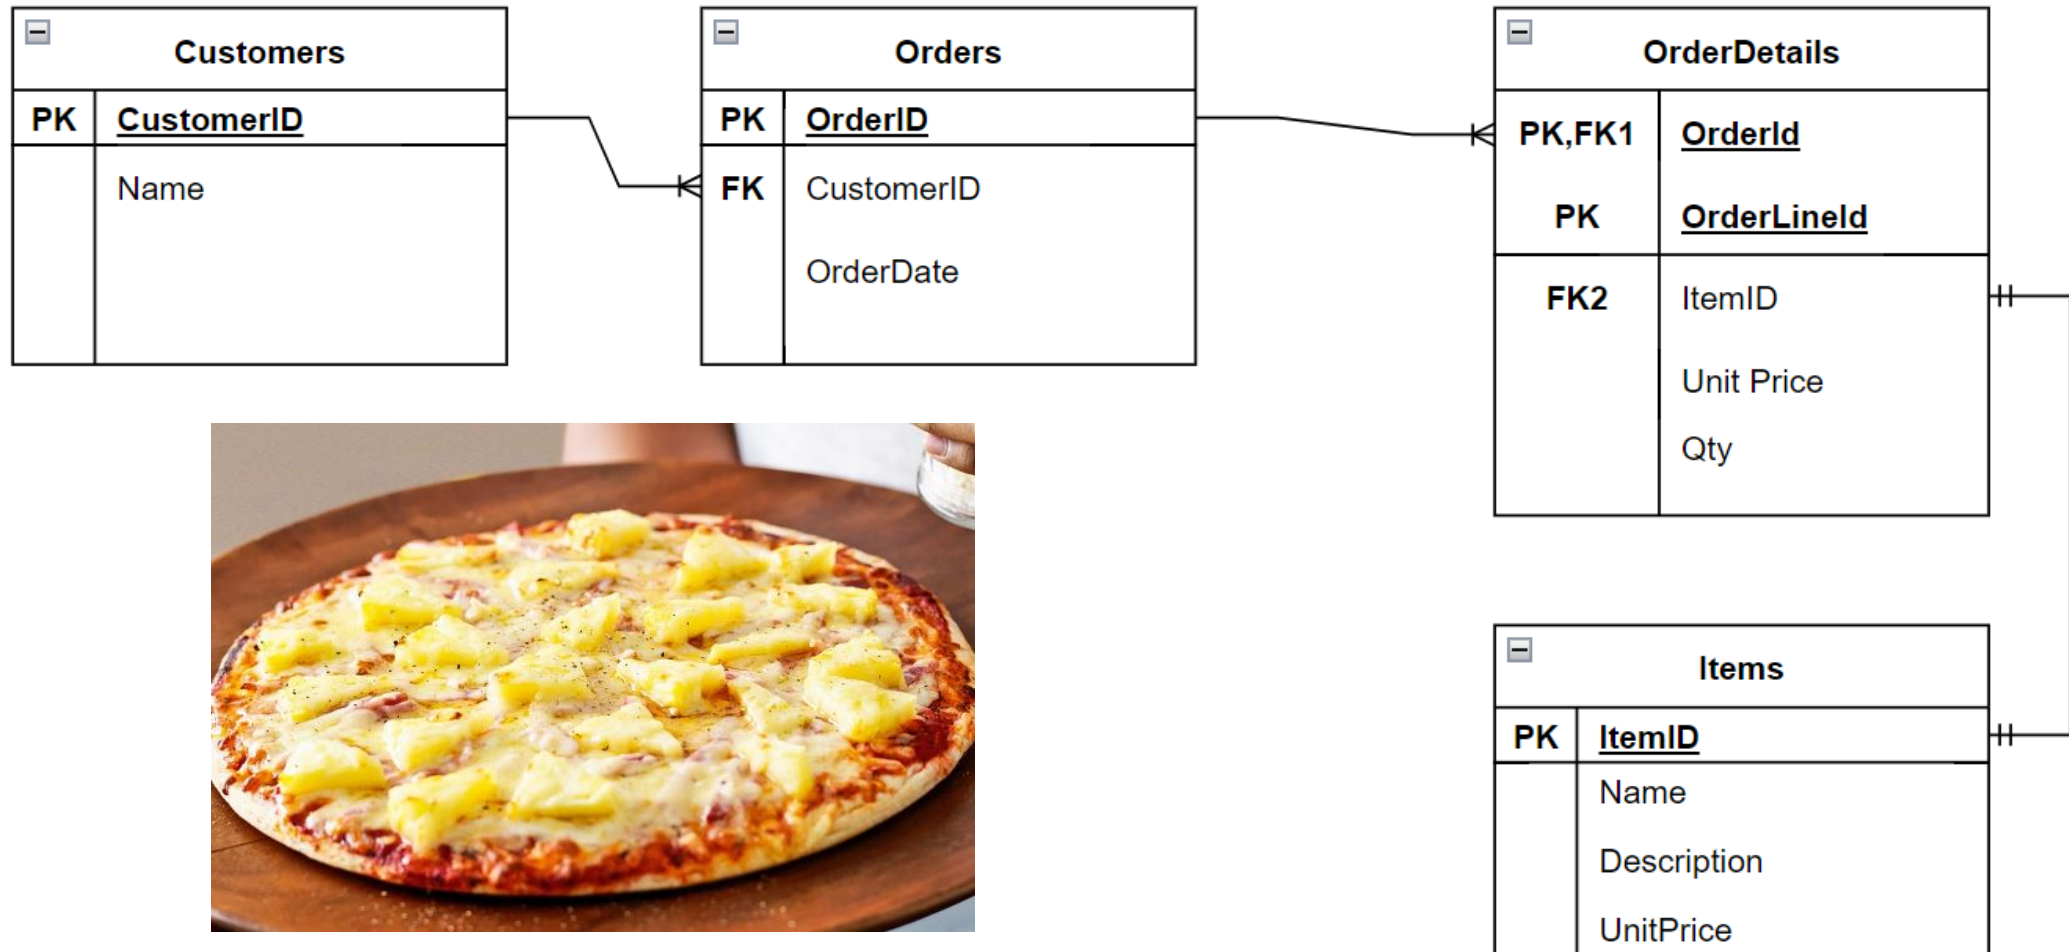




# Data Modeling for Azure Cosmos DB for the Relational Mind

Sarah "sadukie" Dutkiewicz  
Microsoft MVP, Developer Technologies

# Relational Data: Sadukie's Pizza Store



[This Photo](#) by Unknown Author is licensed under [CC BY-NC-ND](#)

How are you  
using this data?



[This Photo](#) by Unknown Author is licensed under [CC BY-SA-NC](#)

# General questions

- ☐ How is this data being used?
- ☐ What fields are queried together?
- ☐ How often does the data change?
- ☐ Is this a read-heavy data set?
- ☐ Is this a write-heavy data set?



[This Photo](#) by Unknown Author is licensed under [CC BY-NC](#)

# Modeling for a Document Database

# Data Modeling for Document Stores

- Schema-less
- How is data being used?
  - Reads?
  - Writes?
- Denormalization
- Embedding vs referencing data
- Partition considerations
- Microsoft Docs - Modeling data in Azure Cosmos DB:  
<https://docs.microsoft.com/en-us/azure/cosmos-db/sql/modeling-data>
- Azure Cosmos DB Essentials Season 2 – Schema Design Strategy:  
<https://www.youtube.com/watch?v=bKDaL-GRSAM>



# Denormalize / Embed

- Contained relationships
- One-to-few relationships
- Infrequent changes
- Queried together
- Data has bounds

```
{
  "CustomerName": "Matt",
  "OrderId": 3,
  "OrderDate": "2022-03-01T16:50:11",
  "OrderDetails": [
    {
      "OrderLineId": 1,
      "ItemName": "Cheese Pizza",
      "Qty": 4,
      "UnitPrice": 10.99
    }
  ]
}
```

# Normalize / Reference

- One-to-many relationships
  - Many-to-many relationships
  - Frequent changes
  - Data has no bounds
- 
- Manually managed – no foreign keys/constraints
- 
- Requires more round trips – no JOINS in the relational SQL sense
  - Typically better for **write-heavy loads**



# How are you using this data?

Demo: Document Store for Order History – Azure Cosmos DB SQL API

Reference:

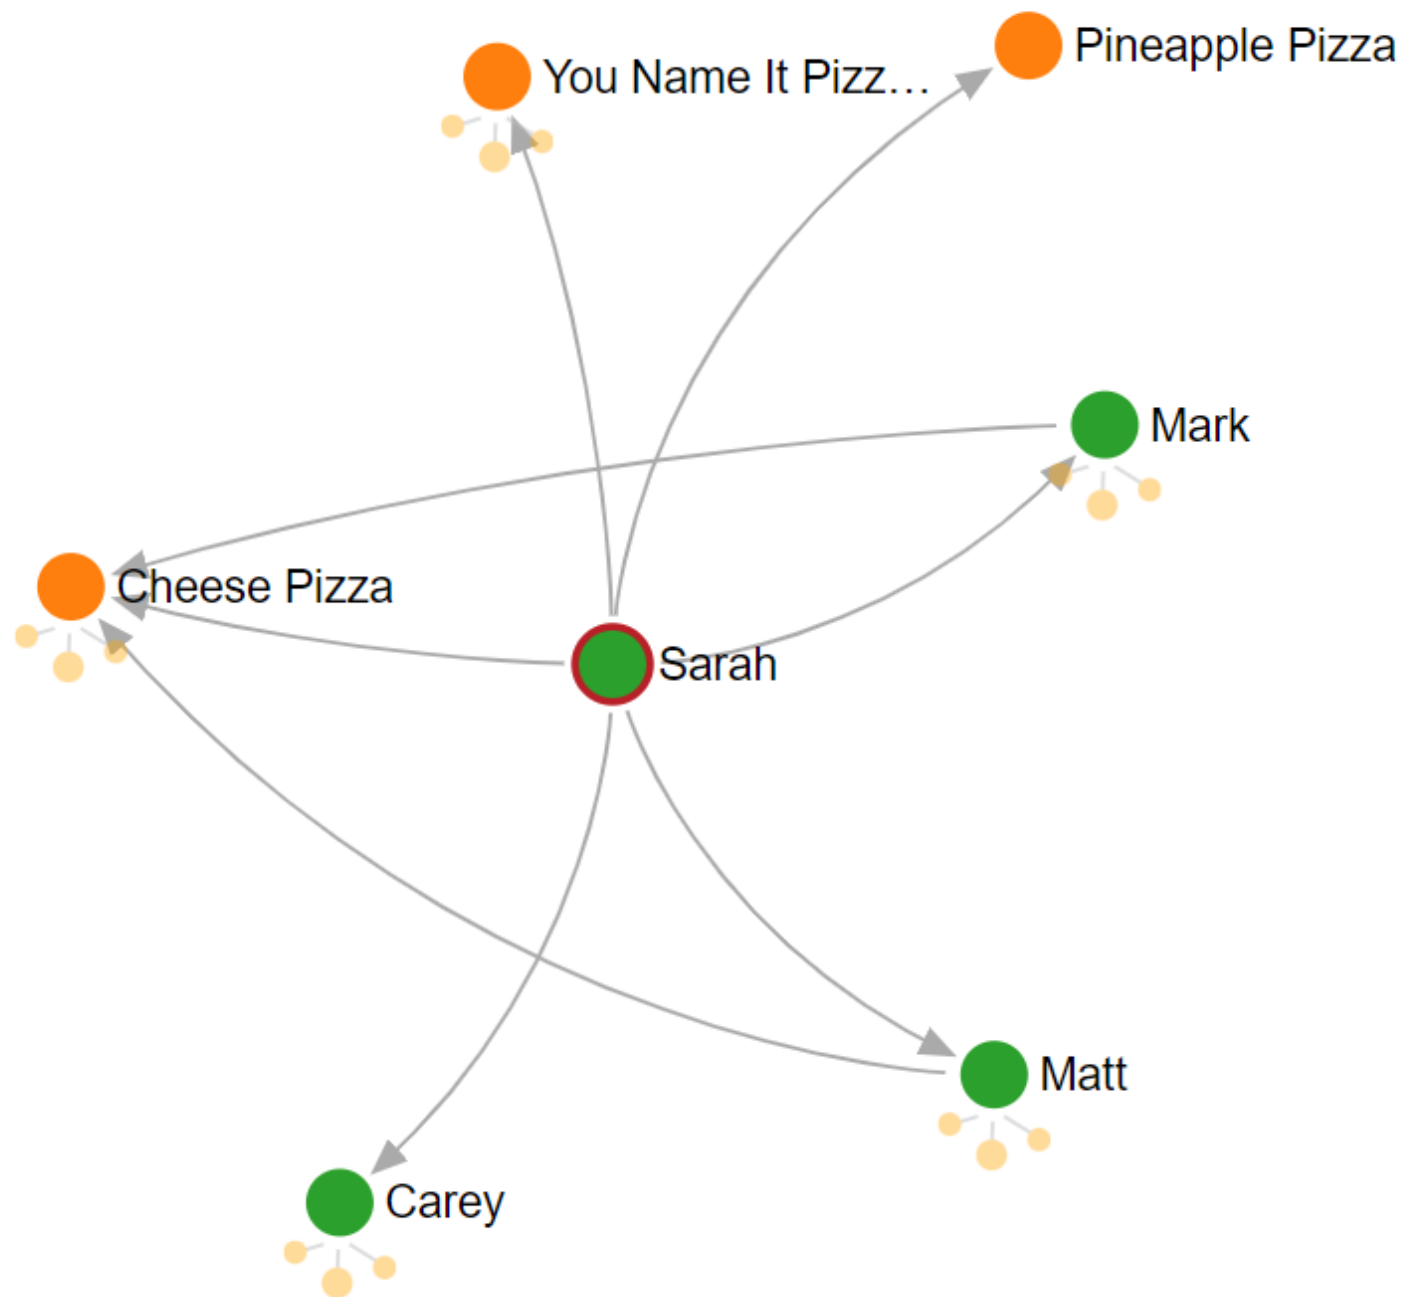
Microsoft Docs – Azure Cosmos DB PDF query cheat sheets:

<https://docs.microsoft.com/en-us/azure/cosmos-db/sql/query-cheat-sheet>



[This Photo](#) by Unknown Author is licensed under [CC BY-SA-NC](#)

# Modeling for a Graph Database



# People who bought pineapple pizza also bought...

- `g.V()`
  - `.has('name','Pineapple Pizza')`
  - `.inE('bought')`
  - `.outV()`
  - `.outE('bought')`
  - `.inV()`
  - `.hasLabel('item')`
- `.executionprofile()`

# Graph Database

- Azure Cosmos DB Gremlin API
- Using Gremlin syntax
- Use while exploring **relationships** (edges) between **things** (vertices)
- Collection of vertices and edges is known as a **graph**
- Common uses:
  - Recommendation engine
  - Fraud detection
  - Managing IT networks – licensing tracking, asset management
  - Software dependency graphs

# JSON under the covers - Vertex

```
▼2: Object
  label: "item"
  id: "e2fc74a4-a108-4c41-b2c5-626ca616cd6c"
  ▼name: Array[1]
    ▼0: Object
      id: "bef80c66-038a-47a6-ae6b-1579a59f7bbc"
      _value: "Pineapple Pizza"
  ▼itemId: Array[1]
    ▼0: Object
      id: "61f92ca0-bd48-45e0-9701-18dcaecdf953"
      _value: 3
  ▼description: Array[1]
    ▼0: Object
      id: "0d69f36a-99d1-4297-a3d9-fd69dc830ed9"
      _value: "For the pineapple lover who can eat it on their pizza"
  type: "item"
```

# JSON under the covers - Edge

```
▼ 5: Object
  label: "referred"
  id: "331873db-3881-4b8b-9d2d-51044befacb8"
  _sink: "25383175-1640-428a-ac5f-5955728b9b98"
  _sinkLabel: "customer"
  _sinkPartition: "customer"
  _vertexId: "717e0df7-5dd6-4eb7-823c-555eeea293b8"
  _vertexLabel: "customer"
  _isEdge: true
  type: "customer"
```

# Graph Database Modeling Considerations

- Identify entities within a single query
  - Entities become vertex objects of a graph.
- Determine direction of relationships
  - Azure Cosmos DB Gremlin API is directional – **in** and **out**
  - Vertices are stored with **out** edges
  - **in** is cross-partition – so if used frequently, add edges in both directions
- Understand **read and write requirements**
  - Helps with graph strategies
- Queries can be evaluated with `.executionprofile()`



# Graphing at Scale

- Azure Cosmos DB horizontal scaling at work
- Partition key is used for balancing
  - Required for Gremlin drivers 3.x
  - Required for Azure portal
- Vertices stored based on partition key
- Edges are stored with their source vertices
- Microsoft Docs – Data partitioning in Azure Cosmos DB GremlinAPI:  
<https://docs.microsoft.com/en-us/azure/cosmos-db/graph/graph-partitioning>

# Graph Strategies

- Vertex-based properties
  - Reduces redundancy – common relational mindset
  - Adds:
    - Model complexity
    - Increased latency
    - Query complexity
    - Computation cost
- Property-embedded vertices
  - Allows properties tied to the entity
  - Key-values
  - Adds redundancy across nodes
  - Reduces model complexity
  - Simpler queries
  - Cost-efficient traversals
- Microsoft Docs – Graph data modeling for Azure Cosmos DB Gremlin API:  
<https://docs.microsoft.com/en-us/azure/cosmos-db/graph/graph-modeling>

# How are you using this data?

Demo: Graph for  
Recommendation Engine –  
Azure Cosmos DB Gremlin  
API

Reference:

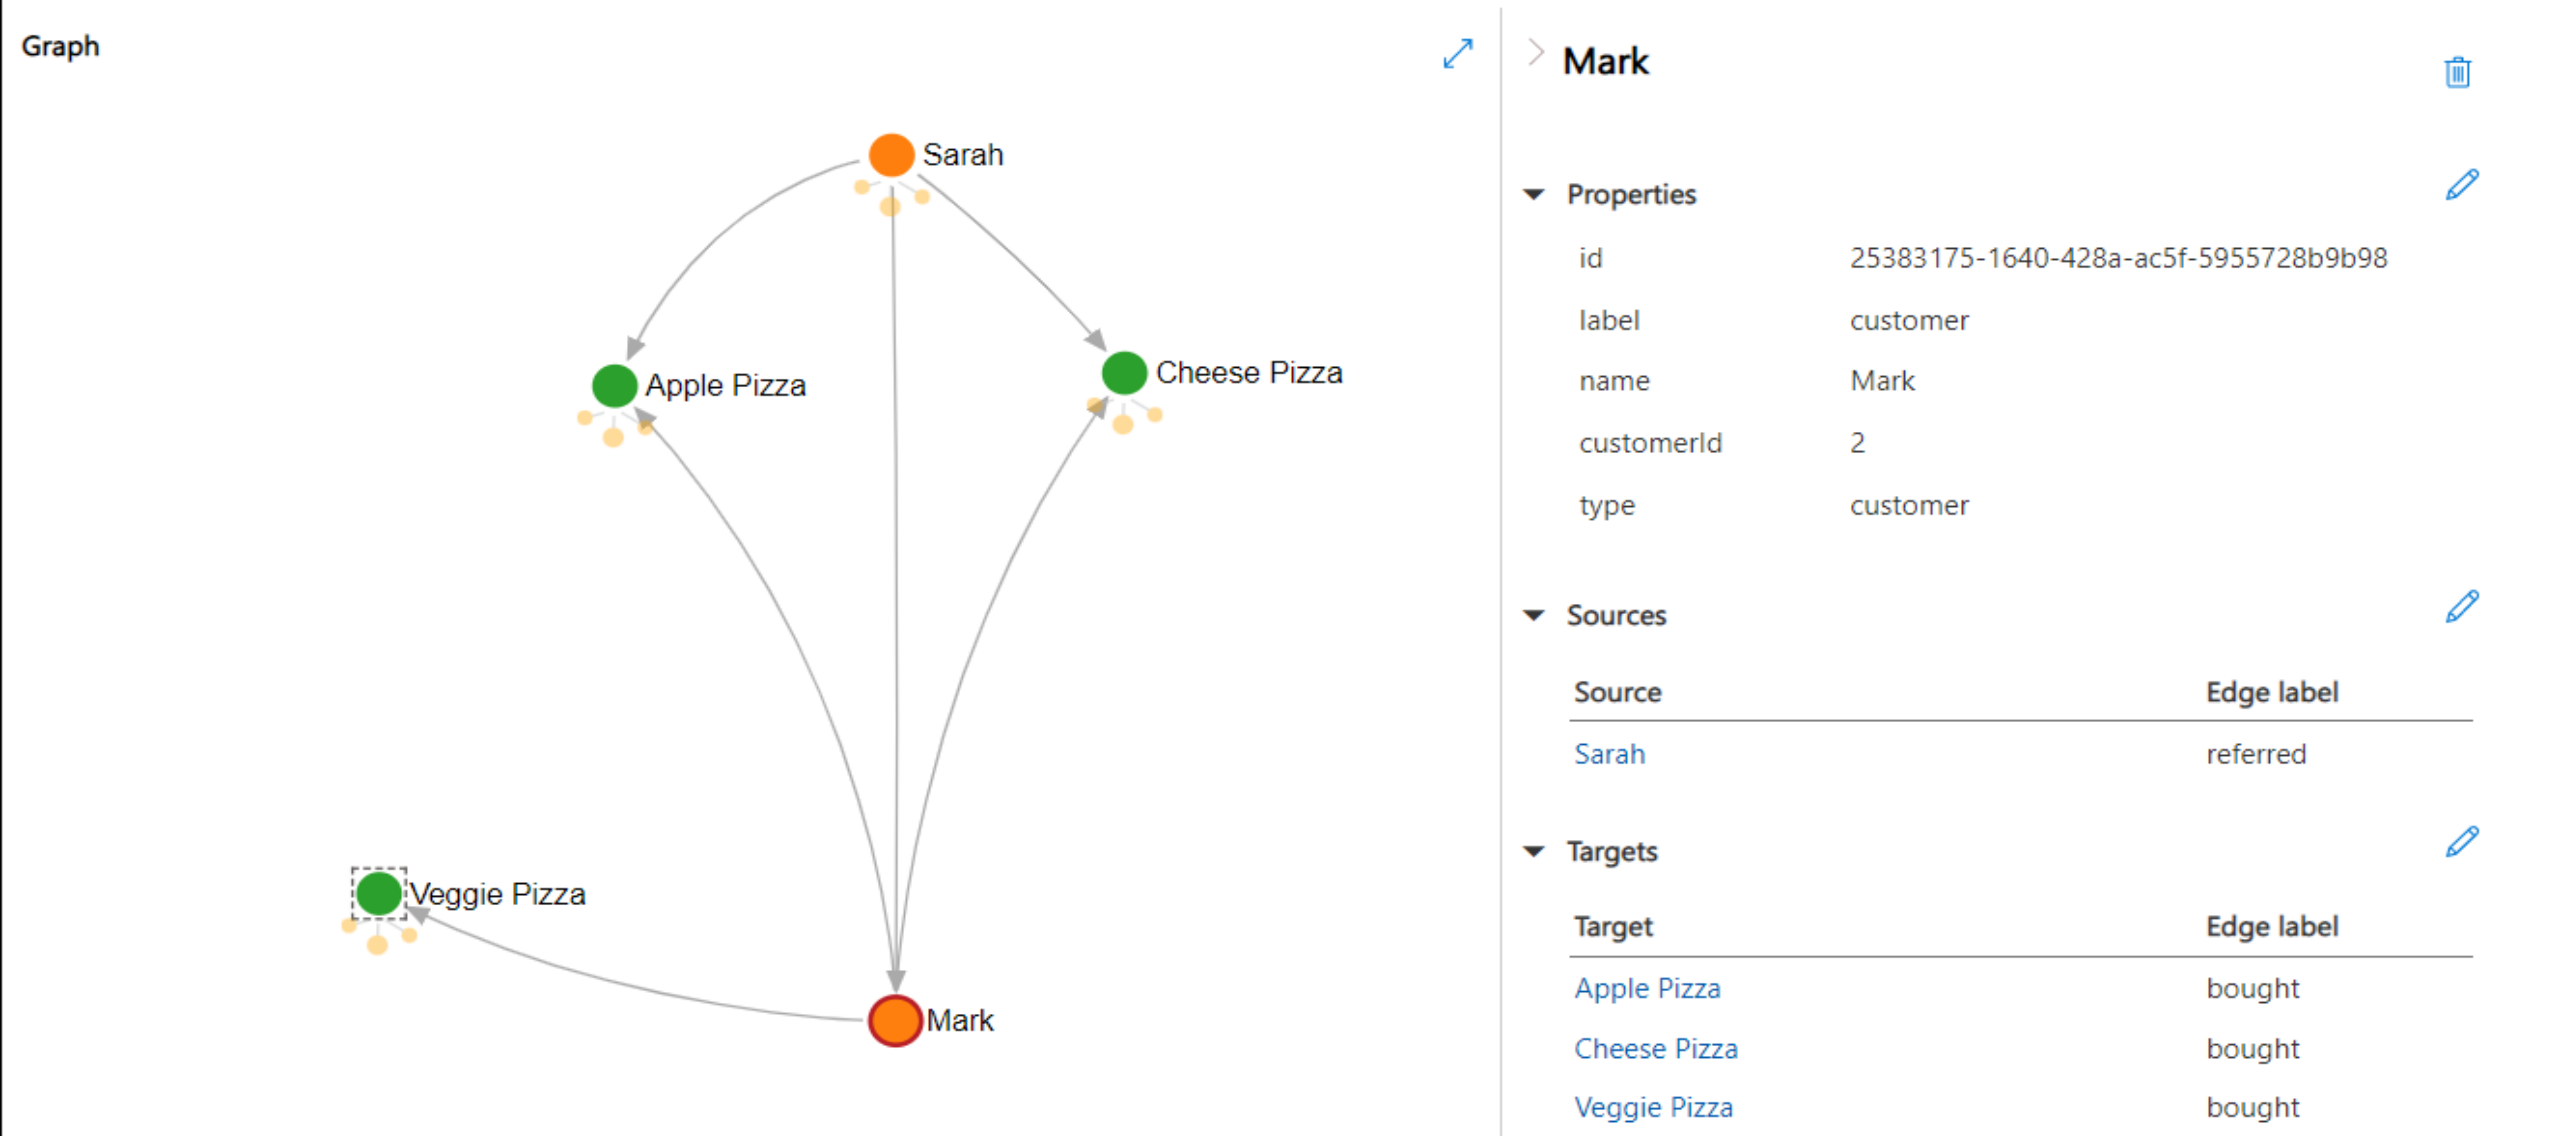
Microsoft Docs – Azure Cosmos DB PDF query  
cheat sheets:

<https://docs.microsoft.com/en-us/azure/cosmos-db/sql/query-cheat-sheet>



[This Photo](#) by Unknown Author is licensed under [CC BY-SA-NC](#)

# Demo – Recommendation Engine



# Resources

# Azure Cosmos DB Resources for Data Modeling

- [Modeling data in Azure Cosmos DB - Azure Cosmos DB | Microsoft Docs](#)
- [Schema Design Strategy | Azure Cosmos DB Essentials Season 2](#)
- [Azure Cosmos DB PDF query cheat sheets | Microsoft Docs](#)
- [Data partitioning in Azure Cosmos DB Gremlin API | Microsoft Docs](#)
- [Graph data modeling for Azure Cosmos DB Gremlin API | Microsoft Docs](#)

# Contact Info

- Twitter: [@sadukie](https://twitter.com/sadukie)
- LinkedIn:  
<https://www.linkedin.com/in/sadukie>

